

A Little Book of Bioinformatics

Leighton Pritchard

2023-01-01

Table of contents

| | |
|--|----|
| Preface to <i>A Little Book of Bioinformatics</i> | 4 |
| 1 Introduction | 5 |
| 2 Summary | 6 |
| I Representing Biological Data | 7 |
| II Biological Databases | 9 |
| III Pairwise Sequence Alignment | 11 |
| IV Multiple Sequence Alignment | 13 |
| 3 MAFFT | 15 |
| 3.1 Where can I download MAFFT? | 15 |
| 3.2 Where can I use MAFFT in a browser? | 15 |
| 3.3 How does MAFFT work? | 15 |
| 3.3.1 Multiple workflows, tuned for different circumstances | 15 |
| 3.3.2 Pairwise sequence, or sequence group alignment | 17 |
| 3.3.3 Scoring scheme | 18 |
| 3.3.4 Guide tree building | 18 |
| 3.3.5 RNA and protein alignment incorporating structural information | 18 |
| 3.3.6 Extension of existing alignments with new sequences | 19 |
| 3.3.7 Parallelisation | 19 |
| 3.4 Where can I find out more about MAFFT? | 19 |
| V Biological Networks | 20 |

| | |
|---|-----------|
| VI Biological Structures | 22 |
| 4 PyMOL | 24 |
| 4.1 Where can I download PyMOL? | 24 |
| 4.2 Using PyMOL | 24 |
| 4.2.1 PyMOL layout | 24 |
| 4.3 Obtaining structures | 26 |
| 4.4 Changing the appearance of the structure | 28 |
| 4.4.1 Rotating the molecule | 28 |
| 4.4.2 Changing molecule colours | 28 |
| 4.4.3 Hiding elements of the structure | 29 |
| 4.4.4 Selecting part of the structure | 29 |
| 5 ChimeraX | 34 |
| 5.1 Where can I download ChimeraX? | 34 |
| 5.2 How do I use ChimeraX? | 34 |
| 5.2.1 ChimeraX layout | 35 |
| 5.3 Loading a structure | 35 |
| 5.4 Changing the appearance of the structure. | 35 |
| 5.4.1 Switching from cartoons to atoms | 35 |
| 5.4.2 Changing atoms representation | 36 |
| 5.4.3 Changing molecule representation | 38 |
| 5.5 Visualising Sequence Conservation | 38 |
| 5.5.1 Requirements | 38 |
| 5.5.2 Loading the alignment data | 39 |
| 5.5.3 Visualising sequence conservation | 40 |
| 5.5.4 Selectively visualising conserved sites | 41 |
| VII Signatures of Evolution | 45 |
| References | 47 |

Preface to *A Little Book of Bioinformatics*

Welcome to *A Little Book of Bioinformatics*. This is an online book, under continual development, which I am building as and when topics come to mind or prominence.

My goal is that this online book will come to be a fairly transparent and honest reference for students in bioinformatics, and maybe for some researchers, too.

I would be very grateful for feedback [by email](#) or through the [GitHub repository Issues page](#)

1 Introduction

This is a book created from markdown and executable code.

See Knuth (1984) for additional discussion of literate programming.

```
1 + 1
```

```
[1] 2
```

2 Summary

In summary, this book has no content whatsoever.

1 + 1

[1] 2

Part I

Representing Biological Data

This section of the book will introduce the computational representation of biological data.

Part II

Biological Databases

This section of the book will introduce biological databases.

Part III

Pairwise Sequence Alignment

This section of the book will introduce the concepts, algorithms, and application of pairwise sequence alignment.

Part IV

Multiple Sequence Alignment

This section of the book will introduce the concepts, algorithms, and application of multiple sequence alignment.

3 MAFFT

MAFFT is a very fast progressive alignment method, suitable for aligning large sequence sets, and tuneable for a range of input circumstances.

MAFFT stands for Multiple Alignment using Fast Fourier Transform, and was first described in 2002 (Katoh et al. 2002). Since then it has received multiple enhancements and extensions to improve scalability to larger sequence sets, and integration of other data such as protein structures (Katoh and Toh 2008a; Katoh and Standley 2013; Rozewicki et al. 2019; Katoh, Rozewicki, and Yamada 2019).

3.1 Where can I download MAFFT?

- <https://MAFFT.cbrc.jp/alignment/software/>

3.2 Where can I use MAFFT in a browser?

- <https://mafft.cbrc.jp/alignment/server/>
- <https://www.ebi.ac.uk/Tools/msa/MAFFT/>
- <https://usegalaxy.eu/>

3.3 How does MAFFT work?

3.3.1 Multiple workflows, tuned for different circumstances

MAFFT can be used as a “black box” alignment tool, but it offers several ways to combine its components. The various combinations can be chosen to suit your alignment task or available hardware, and fall into three main categories:

3.3.1.1 FFT-NS1, FFT-NS-2: progressive methods for large sequence sets

These approaches are simple progressive alignment workflows that take advantage of 6-tuple scoring to speed up initial distance matrix construction, and FFTs to speed up sequence and sequence group alignment.

1. A guide tree is constructed from 6-tuple similarity scores
2. A progressive alignment is performed on the guide tree (and the output alignment returned - FFT-NS-1)
3. The guide tree is reconstructed from this alignment (FFT-NS-2)
4. The sequences are realigned against the rebuilt guide tree (and the output alignment returned - FFT-NS-2)

FFT-NS-2 is the default MAFFT algorithm, chosen to balance speed and accuracy (Katoh and Toh 2008a).

3.3.1.2 FFT-NS-i: progressive method for more accurate alignments of smaller sequence sets

This approach resembles that for FFT-NS-2, but iterates over guide tree reconstruction and sequence realignment until either no improvement can be detected in the alignment (by *Weighted Sum of Pairs*, WSP), or a maximum number of iterations is reached.

1. A guide tree is constructed from 6-tuple similarity scores
2. A progressive alignment is performed on the guide tree
3. The guide tree is reconstructed from this alignment
4. The sequences are realigned against the rebuilt guide tree
5. Steps 3 and 4 are repeated until there is no improvement in the alignment's WSP score, or a maximum number of iterations is reached

3.3.1.3 L-INS-i, E-INS-i, and G-INS-i: consistency scores for difficult alignment cases

These approaches resemble FT-NS-i in that they iterate over guide tree construction and realignment, but differ in that they carry out all-vs-all pairwise alignment initially, rather than the 6-tuple approximation, and employ a COFFEE-like *consistency score* alongside the WSP score, and different alignment algorithms during pairwise alignment.

- E-INS-i uses a global Needleman-Wunsch approach to pairwise alignment, and is suited for alignments between sequences that may have several alignable domains and a high proportion of non-alignable insertions.
- L-INS-i uses a local Smith-Waterman alignment to improve alignments where there is a single alignable domain, with non-alignable flanking sequence.

- G-INS-i uses a global Needleman-Wunsch alignment approach, and is suited to situations where it can be assumed that the entire sequence can be aligned for all inputs.

3.3.2 Pairwise sequence, or sequence group alignment

3.3.2.1 Fast Fourier Transform speed-up

MAFFT improves on the speed of simpler progressive alignment approaches (Feng and Doolittle 1987) by using efficient algorithms to target slow steps in the process. The main advance, which gives MAFFT its name, was to use the [fast Fourier Transform \(FFT\)](#) - an algorithm that converts a signal from one quantitative *domain*, such as time, to a representation in the *frequency domain* - to speed up alignment of protein sequences.

A key insight was to note that the frequency of amino acid substitutions is often determined by whether the substituting residue conserves important physicochemical properties (Katoh et al. 2002). Conservation of such properties is a feature of *neutral* or *near-neutral* substitutions. In the original paper, each amino acid (a) was assigned to a vector of two values, one representing the amino acid residue's volume ($v(a)$), and one representing its polarity ($p(a)$), normalised to zero mean and unit variance over all amino acids.

By converting the amino acid sequences to be aligned to a sequence of these vectors, MAFFT calculates the *correlation* between sequences in terms of their polarity and volume components. In doing this, MAFFT considers a *positional lag* of k sites in the sequence - a parameter controlling the size of sequence region being considered, and the overall correlation is reported for this value, as $c(k)$. By using a fast Fourier Transform for this operation, the CPU time for a set of N input sequences is reduced from $O(N^2)$ to $O(N \log N)$.

This concept can be extended from an alignment between two sequences to an alignment between two already-aligned groups of sequences, by replacing the vector of volume (or polarity) values for a single sequence with a vector that is the linear combination of the volume components for the aligned group (Katoh et al. 2002).

3.3.2.2 Finding homologous sequence segments

If a pair of sequences has homologous regions - conserving polarity and volume - then there will be a corresponding peak in the graph of $c(k)$. This identifies the *positional lag* of the match, but not its location. So, to determine the actual location of the match, a *sliding window analysis* (window size 30) is carried out to identify local homologies for the 20 highest peaks in $c(k)$. A score threshold of 0.7 per site is applied, and any window with a higher value is considered a *homologous segment*. Successive segments identified as *homologous segments* are combined, to a maximum of 150 sites per homologous segment. Segments longer than this are divided into segments of maximum length 150 (Katoh et al. 2002).

3.3.2.3 Aligning sequence pairs

A matrix of homologous segments between two sequences is constructed, and the alignment obtained by the standard *dynamic programming* procedure (Katoh et al. 2002).

3.3.2.4 Application to nucleotide sequences

For nucleotide sequences, rather than using vectors that represent amino acid side chain properties, four-dimensional vectors of A, C, G, T are used at each column. Otherwise the process is identical (Katoh et al. 2002).

3.3.3 Scoring scheme

MAFFT departs from the common approach of using an all-positive-value scoring matrix, instead adopting a normalised similarity matrix. By default, this was originally derived from a 200 PAM log-odds matrix for both protein and nucleotide alignments, and a simplified gap penalty scheme was also employed (Katoh et al. 2002). More recently, the 200 PAM matrix has been replaced as default by BLOSUM62, and the gap penalty scheme was heavily revised (Katoh and Toh 2008a).

3.3.4 Guide tree building

The modified UPGMA tree building method used in the original version of MAFFT did not scale well to very large sequence sets, and was replaced in v6 onwards by the PartTree algorithm (Katoh and Toh 2008a, 2008b). This constructs an approximate tree from unaligned sequences by partitioning the dataset and is $O(N \log N)$ rather than $O(N^2)$, like UPGMA.

3.3.5 RNA and protein alignment incorporating structural information

In MAFFT v6, options were introduced to include structural information when aligning RNA sequences. This takes base pairing probability into account, and applies a novel four-way alignment consistency function (Katoh and Toh 2008a).

MAFFT v7 introduced the ability to inform an alignment by using the ASH structural alignment program and, in particular, the inter-residue distance between alpha carbons for amino acids in the alignment. This is likely to be most useful when aligning sequence-diverse, but structurally similar, proteins (Katoh and Standley 2013).

MAFFT has been integrated with the DASH structural alignment database, to use structural alignment information as a set of constraints on sequence alignment. This integration is available through the [official MAFFT alignment server](#) (Rozewicki et al. 2019).

3.3.6 Extension of existing alignments with new sequences

Further improvements were introduced in MAFFT v7 to enable addition of unaligned sequences into an existing multiple sequence alignment (Katoh and Standley 2013).

3.3.7 Parallelisation

Threading was added to MAFFT v7. Progressive alignment method outputs are identical in non-threaded and threaded runs, but iterative refinement-based alignments may generate different output when threaded (Katoh and Standley 2013).

3.4 Where can I find out more about MAFFT?

- [MAFFT v7 tips](#)
- [MAFFT v7 algorithm](#)

Part V

Biological Networks

This section of the book will introduce network representations in biology.

Part VI

Biological Structures

This section of the book will introduce structural data and visualisation.

4 PyMOL

PyMOL is a widely-used, open source molecular visualisation package, with paid-for (“incentive”) support and updates.

4.1 Where can I download PyMOL?

The official home of PyMOL is at Schrodinger Software, where you can download the “Incentive” version. This requires a licence, which is free for educational use, but not free for academic/commercial/research use.

- <https://pymol.org/2/>

There is an open-source, freely-licensed version of PyMOL, also provided by Schrodinger Software. This can be downloaded from GitHub and self-compiled, or obtained through homebrew:

- <https://github.com/schrodinger/pymol-open-source>
- <https://formulae.brew.sh/formula/pymol>

4.2 Using PyMOL

Start the program from the application icon, or from the command-line (depending on which version you have installed), to obtain the landing screen (Figure 4.1).

4.2.1 PyMOL layout

Pymol is laid out approximately into four quadrants. Clockwise from top left in Figure 4.1 we have:

- The **interactive command window**. At the bottom of this window is the command prompt PyMOL >, where commands to control the visualisation can be typed and executed.
- The next quadrant (top right) contains a set of **function buttons** providing a range of actions for control of the visualisation.

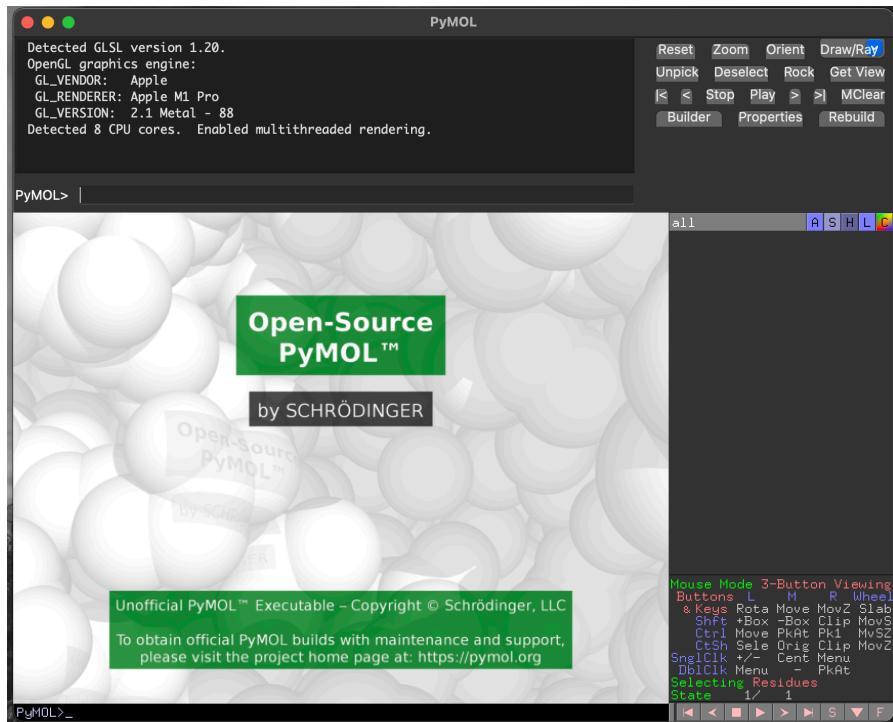


Figure 4.1: PyMOL landing screen, for the Open Source version, on macOS.

- Below this is the **object control panel**, which is the main point-and-click interface for changing the way that molecules in the visualisation window appear, and below this is the mouse control legend to explain what your mouse/trackpad actions will achieve.
- Finally, at bottom left, we have the **main viewer window**, where your molecule will be shown.

4.3 Obtaining structures

PyMOL can connect directly to [RCSB/PDB](#) to download structural data, and can load structure files from your local storage.

4.3.0.1 fetching structures from RCSB/PDB

We will download a structure directly from [RCSB/PDB](#). To do this we will use the **interactive command window**. We also need to know the RCSB/PDB accession for the structure we want to view. Here, we will use the structure [4I61](#), a trimer of PduB, a bacterial microcompartment protein.

 PduB makes biochemical factories inside cells.

PduB is a protein that trimers to form a structural unit that then combines with other similar structural units to enclose a volume within bacterial cells, to make a kind of “factory” for chemical reactions.

The UniProt entry for the protein corresponding to this structure is [F8DQ39](#)

To fetch this structure from RCSB/PDB, we enter the command `fetch 4I61` into the command prompt, and hit `Return` (Figure 4.2).

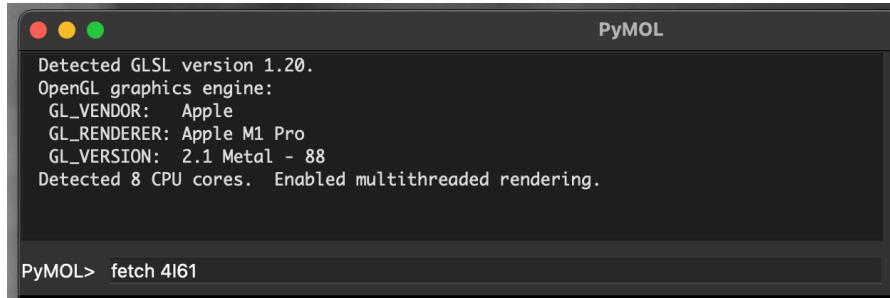


Figure 4.2: PyMOL command prompt including `fetch` command.

Executing this command will produce a short report to the interactive command window, and show a rendering of the structure in the **main viewer window** (Figure 4.3).

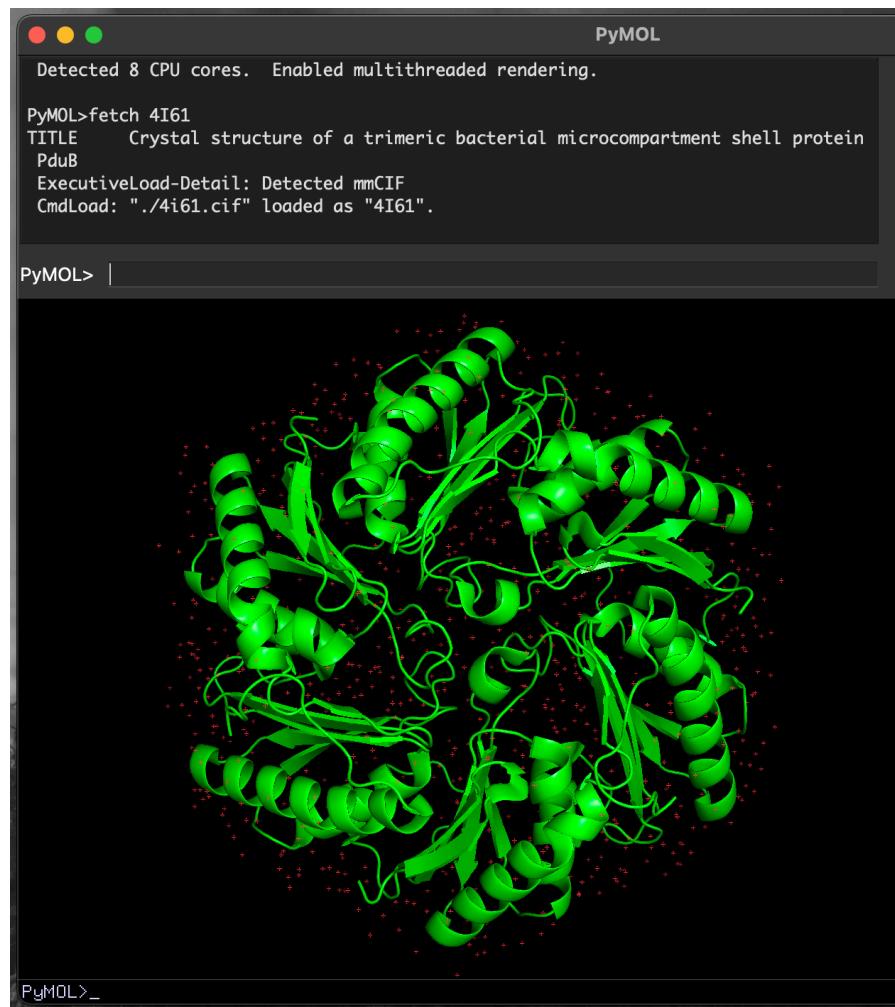


Figure 4.3: Initial render of 4I61 PduB structure in PyMOL.

 PyMOL downloads the fetched file.

When executed, this command will also download the file `4I61.cif` to the current working directory. You can load this file into PyMOL without requiring a live network connection.

4.4 Changing the appearance of the structure

4.4.1 Rotating the molecule

By default, left-clicking on the molecule in the main viewer window and moving your mouse will rotate the molecule. You can use this to obtain a viewing position that helps you understand the structure better or that, when saved as a figure, will communicate your message to a reader.

4.4.2 Changing molecule colours

The **object control panel** provides buttons (A, S, H, L, C) that control aspects of the molecule's appearance:

- A/Action
- S>Show
- H/Hide
- L/Label
- C/Colour

There are three distinct protein chains in this trimeric structure. We will colour the protein differently by chain so that we can see them more clearly.

 Important

- Click on the C button for 4I61 in the **object control panel**
- Click on `by chain` in the menu that appears
- Click on `by chain` in the new menu

This colours each chain in the structure differently, and also the water molecules (small dots) associated with each chain (Figure 4.4).

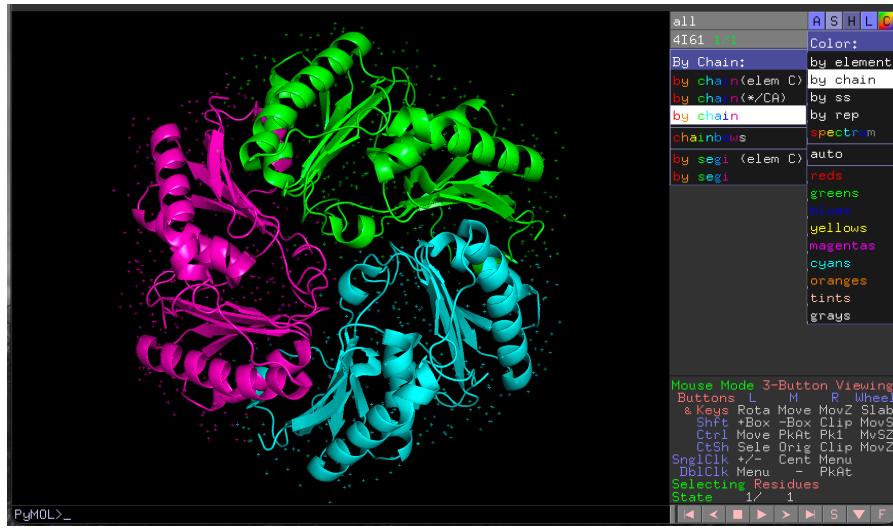


Figure 4.4: PyMOL rendering of 4I61 with a different colour for each chain.

4.4.3 Hiding elements of the structure

We sometimes want to focus attention on particular parts of a structure. To aid in this we can *hide* parts of the visualisation, using the H/Hide menu in the object control panel.

For example, to hide the water molecules that surround each chain, we would:

! Important

- Click on the H button for 4I61 in the object control panel
- Click on **waters**

This removes the water molecules from our visualisation (Figure 4.5).

4.4.4 Selecting part of the structure

It is possible to select parts of the structure in PyMOL by pointing and clicking using the mouse. For complicated selections especially, this can be difficult, tedious, and error-prone. It is usually better to use the **interactive command window** to select structural components explicitly.

4.4.4.1 Selecting with indicate

To select only the first chain of the trimeric PduB structure (chain A), we could execute the command **indicate chain A**, which highlights that chain in the main visualisation window,



Figure 4.5: 4I61 structure with waters removed.

to show that it is selected (Figure 4.6).

See that using `indicate` has produced a new row in the **object control panel** called (`indicate`). This allows us to control the appearance of the selected element.

With the chain selected, we can then change the way it is drawn using the **S/Show** menu for the (`indicate`) selection, and the colours used for rendering, using the corresponding **C/Colour** menu (Figure 4.7, Figure 4.8).

! Important

Click on the (`indicate`) label in the object control panel to cancel the selection.

4.4.4.2 Selecting specific residues with the sequence viewer

Below the mouse action legend on the right hand side of the window there is a set of control buttons. To the right of this list is a button labelled with the letter **S**. Clicking this button brings up the sequence viewer/selector at the top of the main image visualisation window (Figure 4.9).

The scroll bar below the sequence allows you to find any part of the protein sequence (and additional molecules in the structure). The symbols are coloured to match the current structural representation, to aid with locating specific residues. For instance, in Figure 4.10 the residues for chain B are coloured in cyan. Residues can be selected by clicking and dragging within the displayed sequence, and the current selection is highlighted as (`sele`) in the object control panel. We can also select discontinuous sections of the structure, as in Figure 4.11.

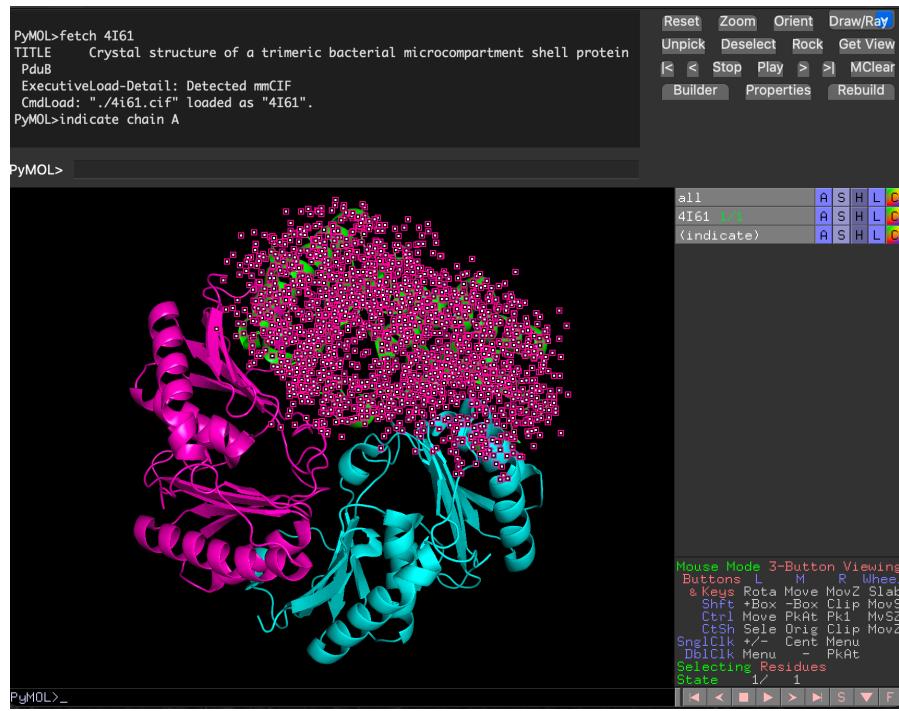


Figure 4.6: 4I61 structure with chain A highlighted.

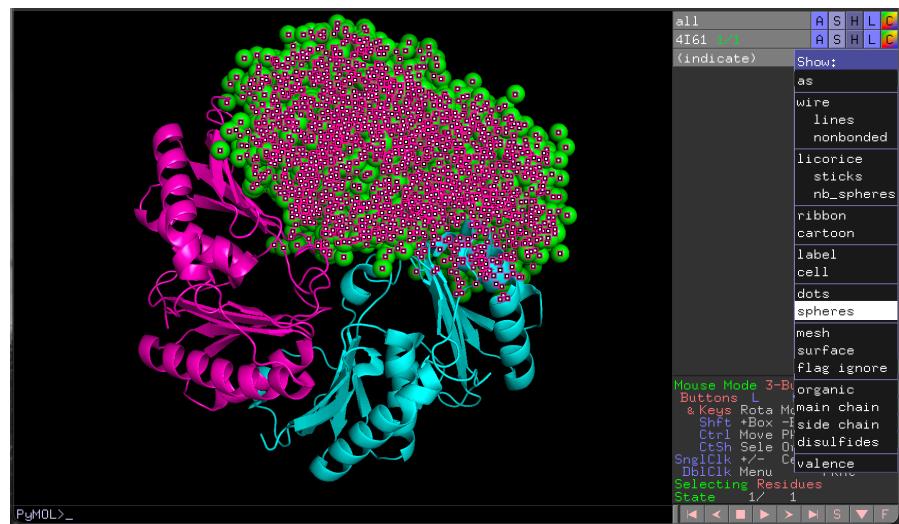


Figure 4.7: Rendering 4I61 chain A structure as spheres.

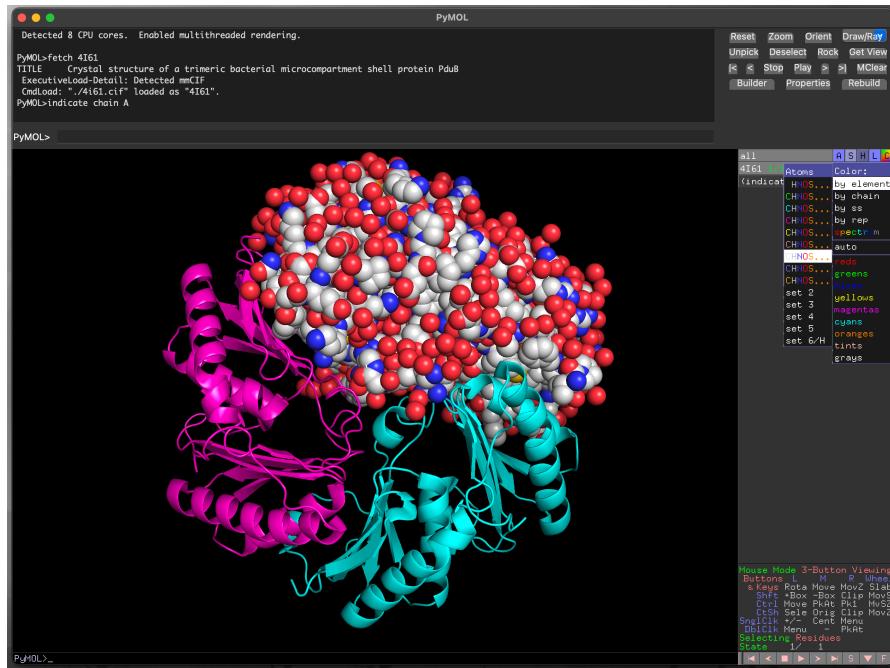


Figure 4.8: Rendering 4I61 chain A structure with atom colouring.

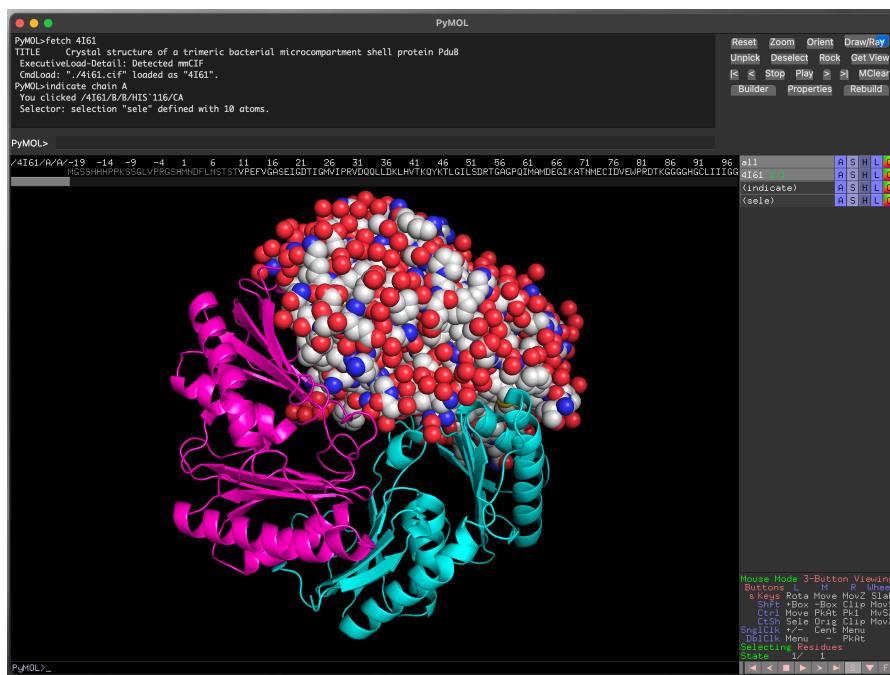


Figure 4.9: Activating the sequence selector.

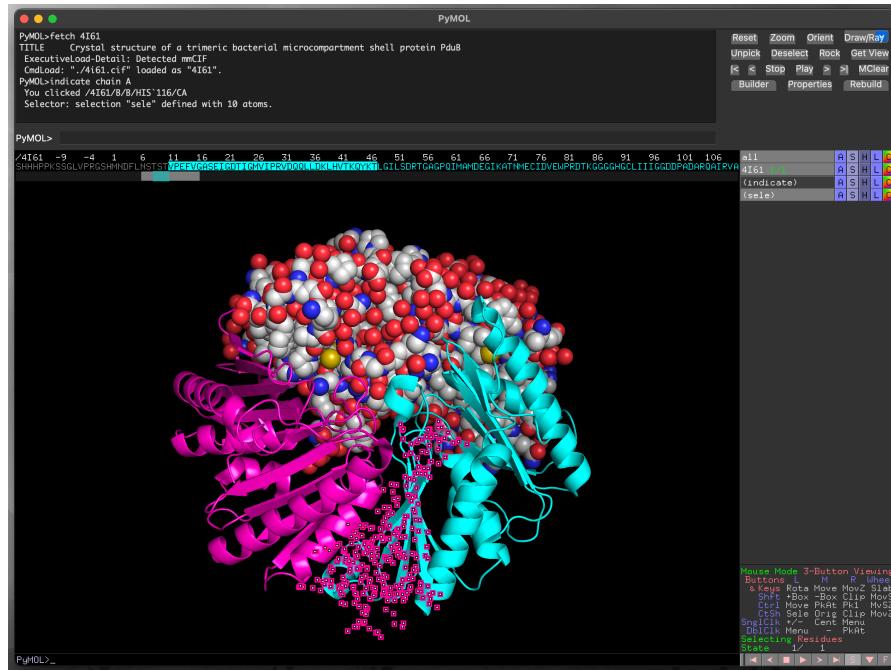


Figure 4.10: Selecting chain residues in PyMOL using the sequence viewer/selector

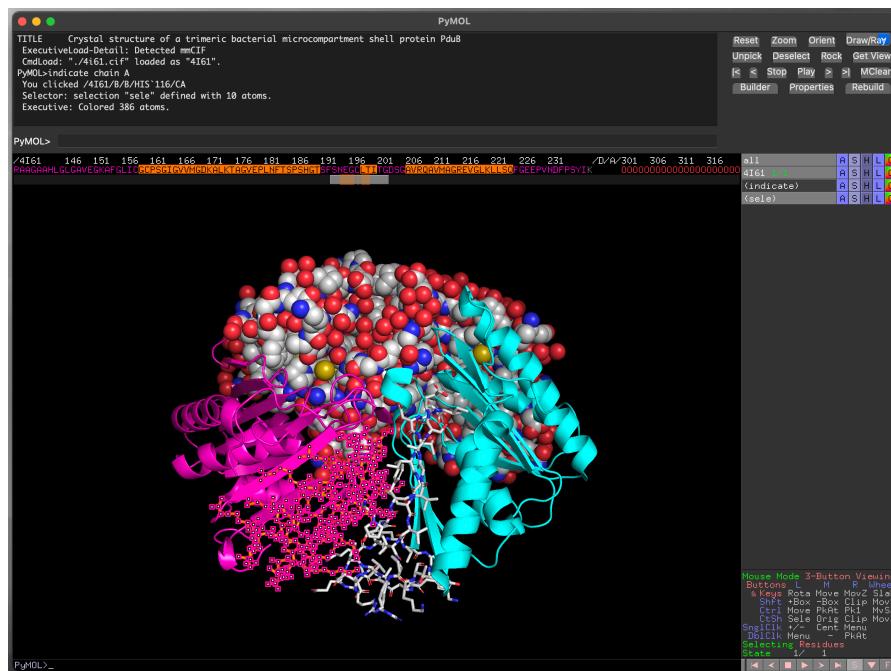


Figure 4.11: Selecting discontinuous residues in PyMOL using the sequence viewer/selector

5 ChimeraX

ChimeraX is a widely-used, open source software package for macromolecular visualisation and analysis. It is free for academic, government, nonprofit, and personal users.

5.1 Where can I download ChimeraX?

ChimeraX is available for download from the UC San Francisco website.

- <https://www.cgl.ucsf.edu/chimerax/download.html>

5.2 How do I use ChimeraX?

Once installed on your machine, use the ChimeraX application icon to start the program. You will be greeted by the standard ChimeraX landing screen (Figure 5.1).

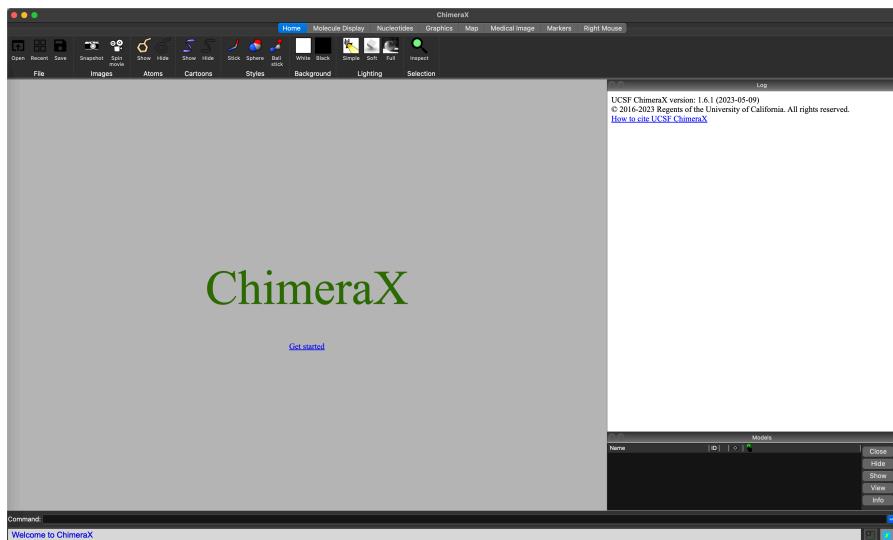


Figure 5.1: ChimeraX landing screen.

5.2.1 ChimeraX layout

The ChimeraX screen is divided into four main areas.

- Along the top of the window is a **ribbon** of icons, and a list of category buttons. Choosing buttons changes the icon/action set.
- Along the bottom of the screen is a **command line interface**, where commands can be entered to control ChimeraX.
- On the right hand side is the **log window** which provides useful information and outputs. This can be toggled on and off with the leftmost button at the bottom right-hand corner (or by entering the `ui windowfill toggle` command in the command line interface).
- The largest part of the screen is the **main visualisation window**, and this can be expanded to fill the full screen width by toggling the log window off.



Tip

The **log window** updates with the equivalent command when icons or buttons are used. This is a useful way to learn commands for scripting ChimeraX.

5.3 Loading a structure

Once you have downloaded a structure from [RCSB/PDB](#), there are three ways to load it into ChimeraX.

1. Click on the **Open** icon in the **Home ribbon**, and use the file dialogue to select your file.
2. Use **File -> Open** in the menu bar (or Cmd-O) to bring up the file dialogue box to select your file.
3. Enter `open [PATH TO FILE]` in the **command line interface**, specifying the file you want to open.

Once the file is loaded, you will be presented with the default view, and information in the **log window** (Figure 5.2). The model will also be visible in the window at the lower right of the screen.

5.4 Changing the appearance of the structure.

5.4.1 Switching from cartoons to atoms

By default, ChimeraX presents a **cartoon view** of the loaded model. In the **Home ribbon**, you can use the **Show and Hide Cartoons** icons to show or hide this representation, and the **Show and Hide Atoms** icons to show or hide an atom-level representation (e.g. Figure 5.3).

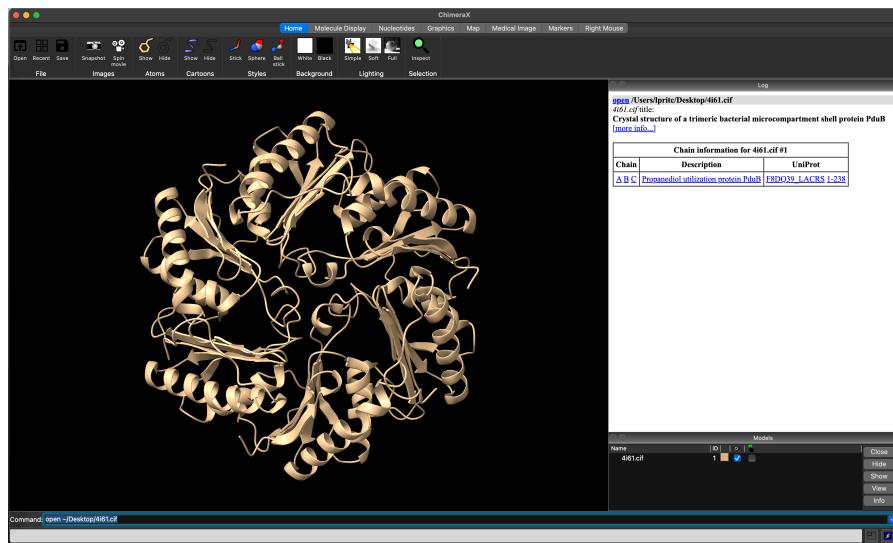


Figure 5.2: ChimeraX view after loading the 4I61 structure.

Change from *cartoon* to *atoms* representation.

1. Click on Show Atoms
2. Click on Hide Cartoons

or in the command line interface, enter

```
show atoms
hide cartoons
```

5.4.2 Changing atoms representation

The Home ribbon provides icons that allow switching between *ball and stick*, *space-filling/sphere* (Figure 5.4), and *stick* representations of the atomic structure.

Command-line equivalents

```
style stick
style sphere
style ball
```

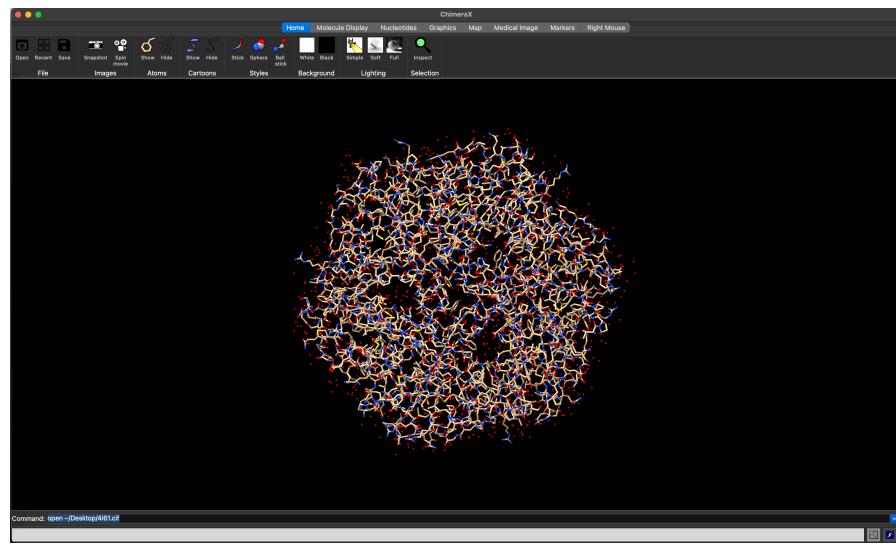


Figure 5.3: ChimeraX atoms view of the 4I61 structure.

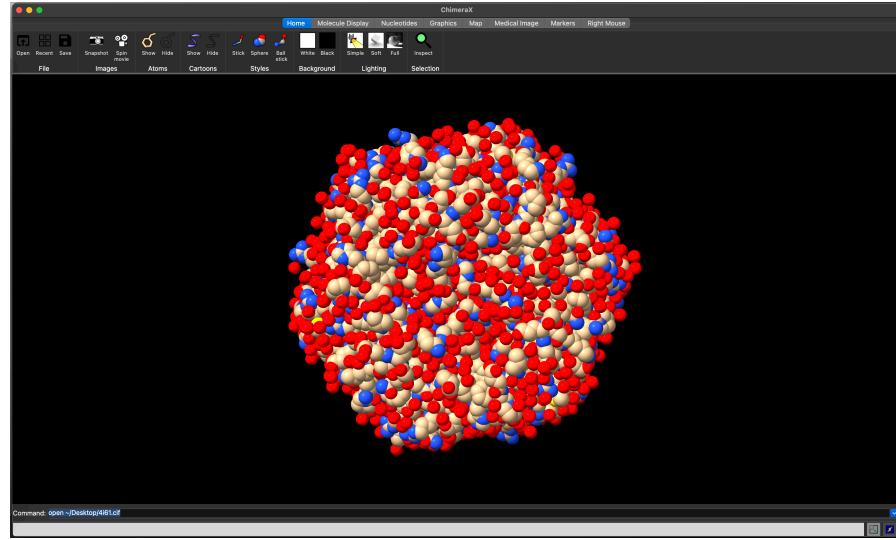


Figure 5.4: ChimeraX space-filling view of the 4I61 structure.

5.4.3 Changing molecule representation

The **Molecule Display ribbon** allows you to control the colouring and other visualisation features for the model. Clicking on the **chain** icon will colour each chain in the model differently, to help visually resolve the overall quaternary structure (Figure 5.5).

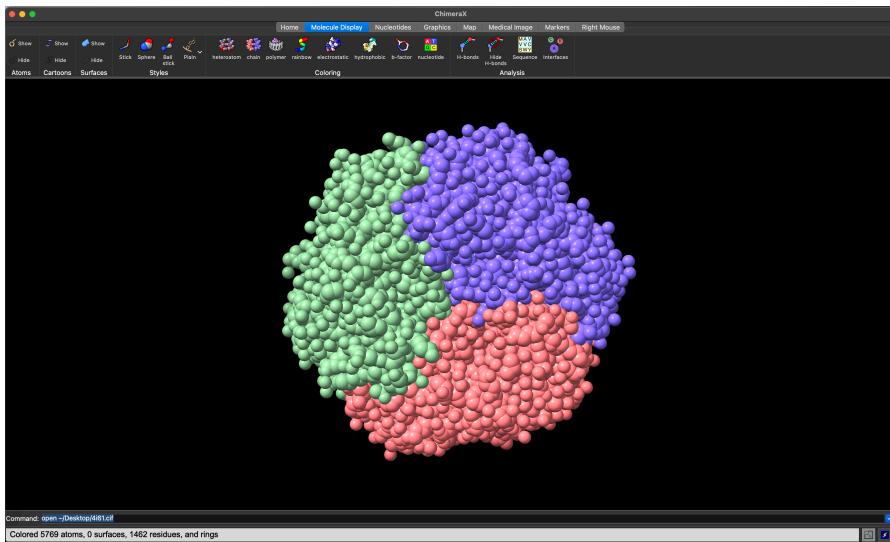


Figure 5.5: ChimeraX space-filling view of the 4I61 structure, with chain colouring.

5.5 Visualising Sequence Conservation

If you have a sequence alignment including the sequence of the protein(s) in your structure, then you can automatically render your model with residues coloured by the extent of sequence conservation.

5.5.1 Requirements

- A structure that includes your protein of interest
- A multiple sequence alignment including your protein of interest

::: :{ .callout-warning } ## Your sequence alignment must meet these criteria

- The sequence corresponding to your protein **must** be the first sequence in the alignment.
- Your alignment **must** be in CLUSTAL format. :::

5.5.2 Loading the alignment data

The sequence alignment can be loaded using the standard **File** → **Open** menu option (Figure 5.6).

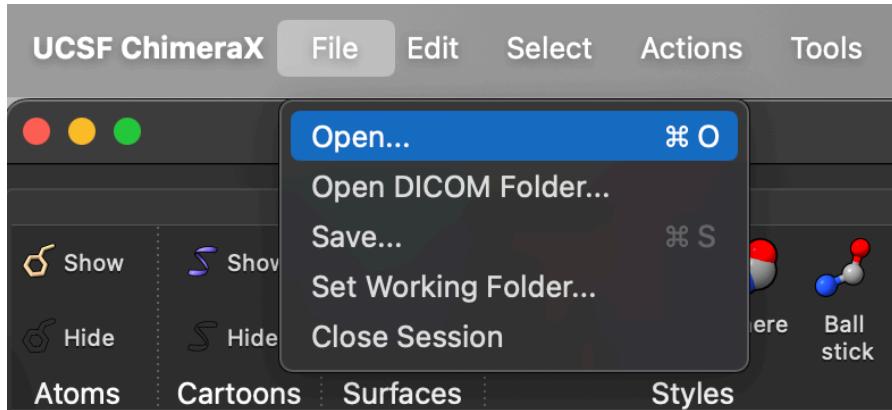


Figure 5.6: ChimeraX file open menu option

Once the file dialogue opens, select the sequence alignment file and click the **Open** button (Figure 5.7).

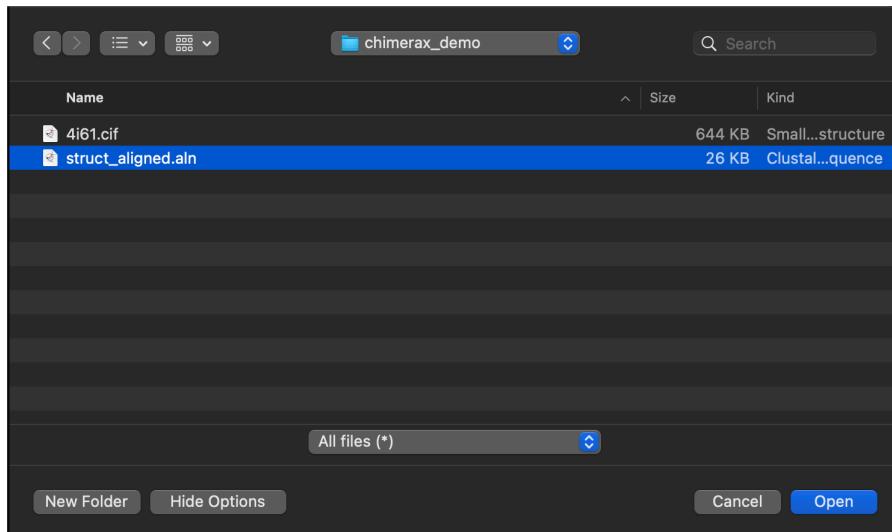


Figure 5.7: ChimeraX file dialogue box. The sequence alignment file is `struct_aligned.aln`

This will open the **Sequence Viewer Window** to display your multiple sequence alignment, with a histogram of sequence conservation at the top (Figure 5.8). The sequence associated with your protein model will be highlighted.

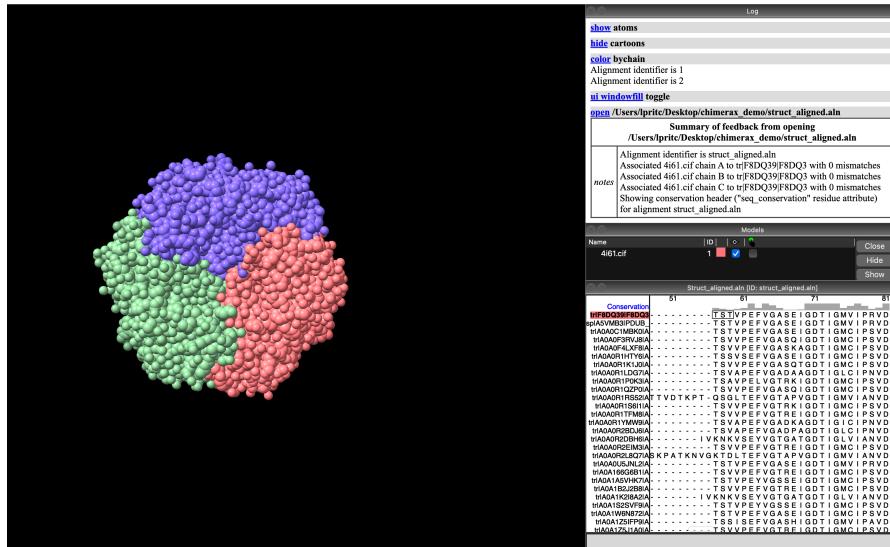


Figure 5.8: ChimeraX with sequence viewer showing alignment for the model’s protein sequences. The sequence associated with the protein model structure (F8DQ39) is highlighted with a red/orange background.

5.5.3 Visualising sequence conservation

ChimeraX associates the sequence conservation, as represented in the histogram in the **Sequence Viewer Window**, with a *variable* called `seq_conservation`. ChimeraX can use this to change the representation of the protein model by assigning the variable to some attribute, such as `color` (for rendered colour).

We can change this attribute using the **command line interface**, by entering the command `color byattr seq_conservation` (Figure 5.9).

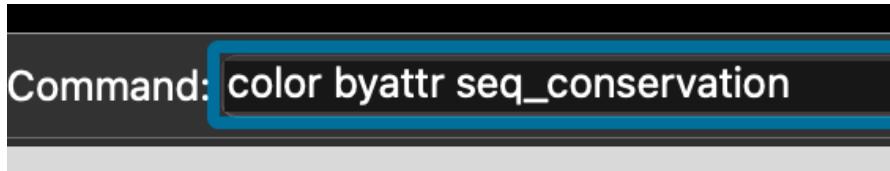


Figure 5.9: ChimeraX command to colour residues by sequence conservation (As defined in a pre-loaded alignment).

This results, by default, in colouring *conserved* sites in red (the more intense the hue, the more conserved the site), and the *variable* sites in blue (again, the more intense the hue, the more variable the site). It can be helpful to change the model representation from space-filling *sphere* to *cartoon* or similar, so that the interior of the protein can be seen more clearly (Figure 5.10).

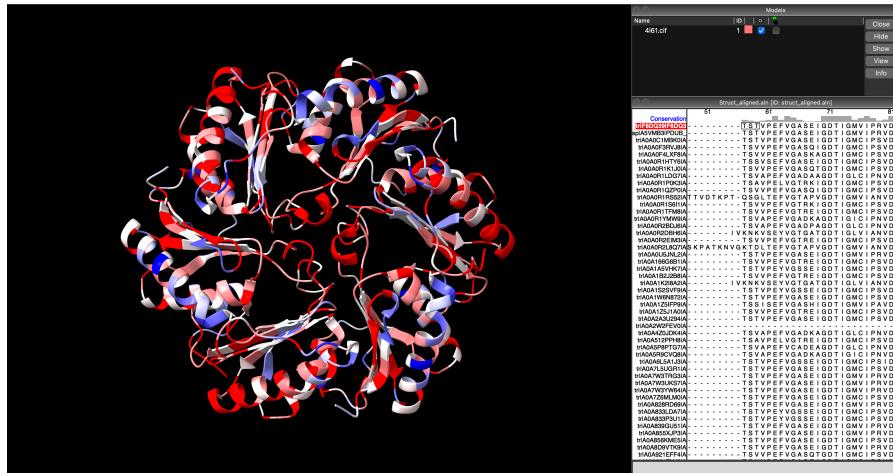


Figure 5.10: ChimeraX cartoon visualisation of 4I61 showing highly-conserved sites in red, and variable sites in blue. Note how the conserved sites are located where the three subunits meet, and where the hexagonal trimer would interact with other trimers (suggesting selection pressure to retain a particular interaction and conformation), and how the variable sites are primarily located on the convex, outward-facing face of the hexagonal subunit.

5.5.4 Selectively visualising conserved sites

We can select sites in the model by attributes, such as sequence conservation, using the `select` command. For example, to select all sites with a conservation score above 0.5 (i.e. the more conserved sites) we would use `select ::seq_conservation > 0.5` (Figure 5.11).

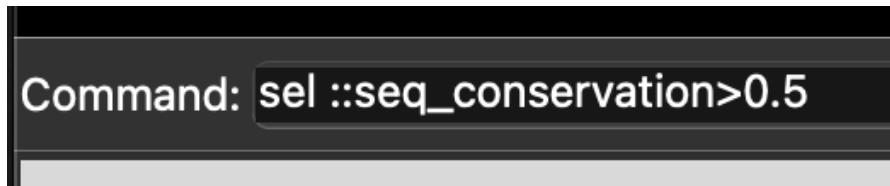


Figure 5.11: ChimeraX command to select residues above a 0.5 sequence conservation score.

Executing this command will select only the sites with a sequence conservation score above 0.5, and highlight them on the current visualisation with a green outline (Figure 5.12).

With only these sites selected, we can show the atom representations for only these sites by clicking on the `Show Atoms` icon, or using the command `show sel atoms`. In (Figure 5.13) this represents those sites as space-filling spheres.

We can then colour the model by sequence conservation as before with `color byattr seq_conservation`, then select only the *low conservation* sites with the command `select`

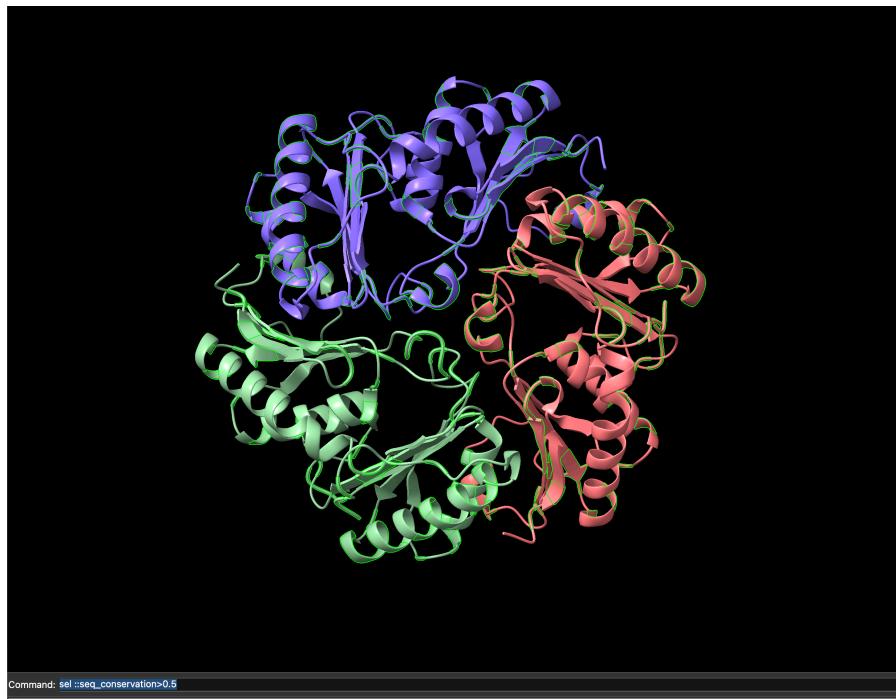


Figure 5.12: ChimeraX visualisation with sites having above 0.5 sequence conservation score highlighted in green.

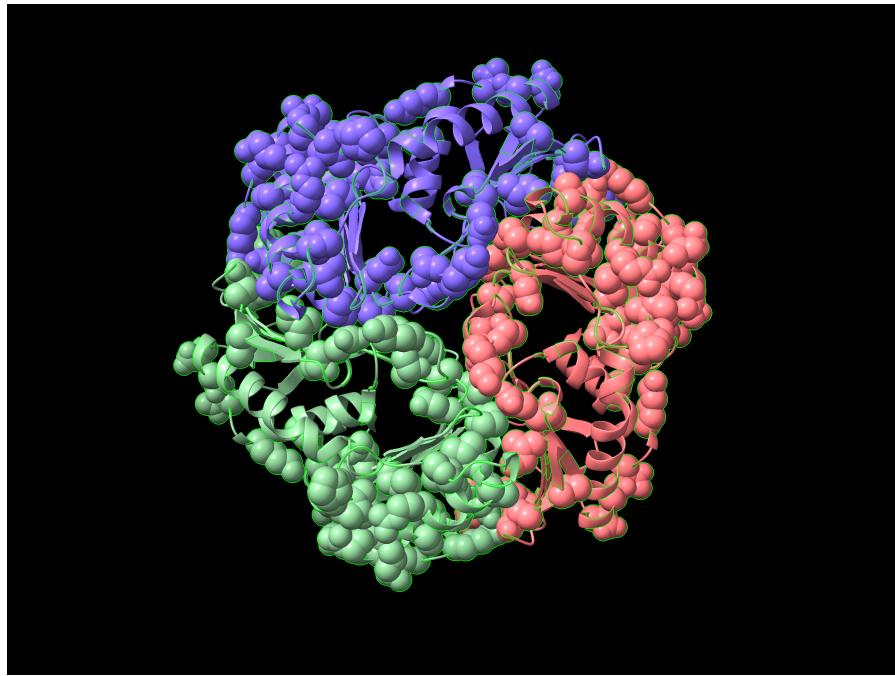


Figure 5.13: ChimeraX visualisation with sites having above 0.5 sequence conservation score rendered as space-filling spheres.

`::seq_conservation < 0.5`, and recolour them by chain (`color sel bychain`), to obtain the image in Figure 5.14, where the protein chains are coloured differently, but highly-conserved residues are indicated in red and space-filling form.

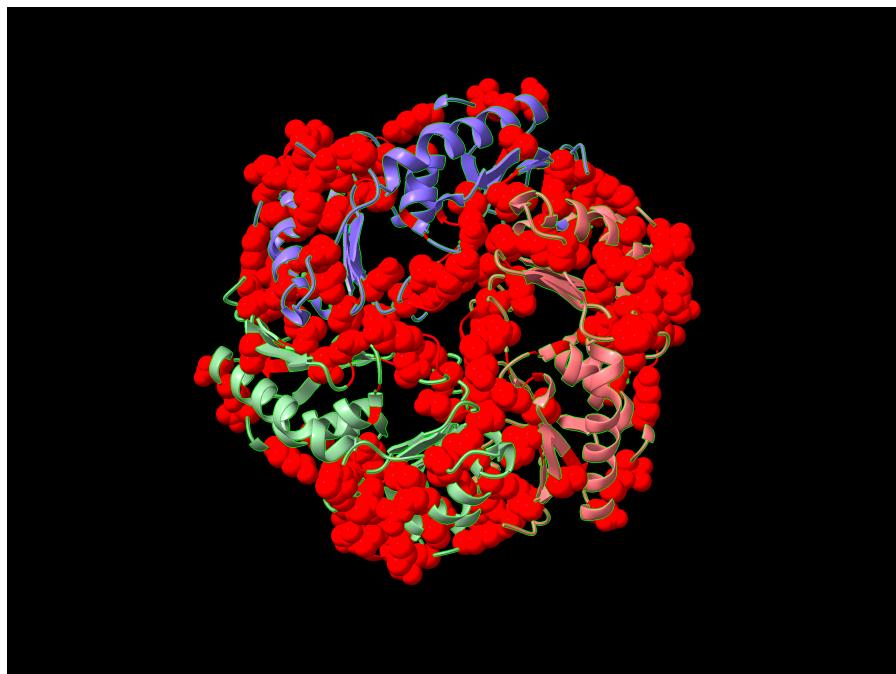


Figure 5.14: ChimeraX visualisation with sites having above 0.5 sequence conservation score rendered as space-filling spheres and coloured by sequence conservation.

Part VII

Signatures of Evolution

This section of the book will introduce phylogenetics and related approaches.

References

- Feng, Da-Fei, and Russell F Doolittle. 1987. "Progressive Sequence Alignment as a Prerequisite to Correct Phylogenetic Trees." *J. Mol. Evol.* 25 (4): 351–60. <https://doi.org/10.1007/BF02603120>.
- Katoh, Kazutaka, Kazuharu Misawa, Kei-Ichi Kuma, and Takashi Miyata. 2002. "MAFFT: A Novel Method for Rapid Multiple Sequence Alignment Based on Fast Fourier Transform." *Nucleic Acids Res.* 30 (14): 3059–66. <https://doi.org/10.1093/nar/gkf436>.
- Katoh, Kazutaka, John Rozewicki, and Kazunori D Yamada. 2019. "MAFFT Online Service: Multiple Sequence Alignment, Interactive Sequence Choice and Visualization." *Brief. Bioinform.* 20 (4): 1160–66. <https://doi.org/10.1093/bib/bbx108>.
- Katoh, Kazutaka, and Daron M Standley. 2013. "MAFFT Multiple Sequence Alignment Software Version 7: Improvements in Performance and Usability." *Mol. Biol. Evol.* 30 (4): 772–80. <https://doi.org/10.1093/molbev/mst010>.
- Katoh, Kazutaka, and Hiroyuki Toh. 2008a. "Recent Developments in the MAFFT Multiple Sequence Alignment Program." *Brief. Bioinform.* 9 (4): 286–98. <https://doi.org/10.1093/bib/bbn013>.
- . 2008b. "Recent Developments in the MAFFT Multiple Sequence Alignment Program." *Brief. Bioinform.* 9 (4): 286–98. <https://doi.org/10.1093/bib/bbn013>.
- Knuth, Donald E. 1984. "Literate Programming." *Comput. J.* 27 (2): 97–111. <https://doi.org/10.1093/comjnl/27.2.97>.
- Rozewicki, John, Songling Li, Karlou Mar Amada, Daron M Standley, and Kazutaka Katoh. 2019. "MAFFT-DASH: Integrated Protein Sequence and Structural Alignment." *Nucleic Acids Res.* 47 (W1): W5–10. <https://doi.org/10.1093/nar/gkz342>.